

**A role-aware explainable AI dashboard for at-risk student
classification: Design and technical evaluation on the Open
University Learning Analytics Dataset**

Daniel Larry^{a,*}, Zainab Shamsiya Usman^a

^a *Department of Computer Science, School of Information Technology and Computing, American
University of Nigeria, Yola, Adamawa State, Nigeria*

* Corresponding author. E-mail address: daniel.larry@aun.edu.ng

Running title: Role-aware explainable AI dashboard for at-risk students

10 **Highlights**

- 11 • One prediction system serves instructors, advisors and administrators separately.
- 12 • XGBoost reached 0.909 recall and 0.887 F1 with course presentations held out.
- 13 • SHAP and LIME shared 7 of 10 top features but only 0.635 top-5 overlap per student.
- 14 • Accuracy varied by over seven points across the four held-out presentations.
- 15 • CommonKADS traces each dashboard view to a stated role-specific knowledge need.
- 16

17 **Abstract**

18 Predictive models built on learning management system (LMS) data can flag students at risk
19 of failing or withdrawing, but most systems present isolated risk scores without the reasoning
20 behind them and give every staff role the same view. This study designed, implemented and
21 technically evaluated a role-aware explainable artificial intelligence (XAI) dashboard that
22 delivers explanations from one prediction system to instructors, academic advisors and
23 administrators at different levels of aggregation. Following design science research with the
24 CommonKADS knowledge engineering methodology, Random Forest, XGBoost and
25 Gradient Boosting classifiers were evaluated on the Open University Learning Analytics
26 Dataset (32,593 student-module-presentation observations) using four-fold cross-validation
27 grouped by course presentation. SHAP and LIME formed the explanation layer, and the
28 prototype was integrated with Moodle and Canvas. The XGBoost prototype, selected on
29 recall, achieved out-of-fold accuracy of 0.877, precision of 0.866, recall of 0.909, F1-score of
30 0.887 and AUC of 0.962; Random Forest obtained a marginally higher F1-score (0.898).
31 Accuracy ranged from 0.846 to 0.916 across held-out presentations. Both explanation
32 methods ranked the number of assessments submitted first and shared seven of their top ten
33 features, but the mean instance-level top-5 overlap was only 0.635. A role-aware explanation
34 layer is technically feasible at full dataset scale, and the partial agreement between SHAP and
35 LIME supports showing both methods side by side. The features cover the full presentation,
36 so results describe retrospective classification; practitioner trust, calibration and fairness
37 remain to be evaluated.

38 **Keywords:** Explainable artificial intelligence; Learning analytics dashboards; At-risk student
39 prediction; SHAP; LIME; Knowledge engineering; OULAD

40

41 **1. Introduction**

42 Learning management systems (LMS) such as Moodle, Canvas and Blackboard now support
43 course delivery, assessment and communication at scale, and their use grew substantially
44 when teaching moved online during the COVID-19 pandemic, increasing the volume of
45 behavioural data available for analysis (Rebelo Marcolino et al., 2025). Learning analytics,
46 defined as the measurement, collection, analysis and reporting of data about learners and their
47 contexts to understand and optimise learning (Siemens & Long, 2011), has used these data to
48 identify students at risk of failing or withdrawing early enough for staff to intervene (Jin et
49 al., 2024; Kuzilek et al., 2017). The Open University Learning Analytics Dataset (OULAD),
50 which records 32,593 student-module-presentation observations across 22 module
51 presentations (Kuzilek et al., 2017), has become one of the most widely used benchmarks for
52 this work (Gunasekara & Saarela, 2025; Jin et al., 2024; Wang, 2025).

53 Predictive accuracy on such benchmarks is now high, but less attention has been given
54 to what happens after a student is flagged. Many prediction systems generate risk scores
55 without explaining the reasoning behind them, leaving staff without enough context to decide
56 whether and how to act (Susnjak et al., 2022). Explainable AI (XAI) techniques such as
57 SHapley Additive exPlanations (SHAP) and Local Interpretable Model-agnostic Explanations
58 (LIME) can attribute a prediction to the features that drove it, yet their integration into
59 working learning analytics dashboards remains limited (Ahmed et al., 2025; Ali et al., 2023).
60 A further problem is that institutional stakeholders need different kinds of explanation from
61 the same model: an instructor monitoring a running course does not need the same view as an
62 administrator reviewing withdrawal across a cohort (Kostopoulos et al., 2024). The literature
63 suggests that explanations should reflect the recipient's role and expertise (Haque et al.,
64 2023; Kostopoulos et al., 2024), but the reviewed studies did not identify a dashboard that
65 delivers role-specific explanations to several staff roles from a single prediction system
66 (Afzaal et al., 2021; Cabral et al., 2025; Kaliisa et al., 2023; Susnjak et al., 2022).

67 This study treats the problem as both a design problem and a technical one. It designs,
68 implements and technically evaluates a role-aware XAI dashboard for at-risk student

classification, guided by the CommonKADS knowledge engineering methodology (Schreiber et al., 2000), and asks three research questions:

- RQ1. How can SHAP and LIME be integrated into a knowledge engineering pipeline to improve transparency in student at-risk prediction models?
- RQ2. What dashboard features and explanation structures, based on the CommonKADS Agent Model and the learning analytics dashboard literature, can be specified to support the different information needs of instructors, advisors and administrators within an LMS?
- RQ3. What do the predictive performance and explainability findings indicate about the technical potential of the role-aware dashboard to support practitioner interpretation and decision-making?

The study makes four contributions: (1) an implemented role-aware XAI dashboard architecture that differentiates SHAP and LIME explanations across instructor, advisor and administrator views of one LMS-integrated prediction system; (2) a documented ensemble model evaluated on OULAD under a protocol that holds out entire course presentations; (3) technical evidence on the agreement and divergence of SHAP and LIME on a real educational dataset; and (4) candidate design guidelines for explainable learning analytics dashboards, derived from the CommonKADS models and the implemented system. The evaluation is technical and uses secondary data only. Findings are therefore reported at three levels of evidence: properties demonstrated by measurement, properties supported by the design but not tested with users, and properties not empirically established.

2. Related work

2.1. At-risk prediction and evaluation protocols

Research on at-risk prediction spans logistic regression, decision trees, ensembles and deep learning. In a systematic review of OULAD studies, Jin et al. (2024) found that ensemble methods, particularly gradient boosting variants and random forests, are common and typically outperform simpler classifiers, while practitioner-facing aspects receive little attention. Reported accuracies can be very high: Wang (2025) reports 96.9% accuracy on

OULAD for a hybrid ensemble, and Rebelo Marcolino et al. (2025) report strong dropout prediction on Moodle logs using CatBoost with multi-objective optimisation. Guevara-Reyes et al. (2025) show that more complex ensembles gain accuracy at the cost of intrinsic interpretability. Engagement-based features tend to outperform static demographic ones (Hussain et al., 2018), useful predictions can be made after a fraction of a course has elapsed (Adnan et al., 2021), and deadline-anchored activity windows have been reported to give stronger early-warning signals on OULAD than calendar windows (Hlosta et al., 2017).

Reported figures depend strongly on the evaluation protocol. Jin et al. (2024) note that protocols vary considerably across OULAD studies, Krstajic et al. (2014) show that letting observations from the same higher-level unit appear in both training and validation folds produces optimistic estimates, and Pelánek (2018) argues that cross-validation, filtering and metric choices materially change reported performance in student modelling independently of model quality. Fairness is a further concern: Baker and Hawn (2022) document systematic performance disparities across demographic subgroups in educational models, and Yu et al. (2020) show that error rates can differ across student groups even when overall accuracy is high.

2.2. Explainable AI in education

XAI surveys distinguish intrinsically interpretable models from post-hoc explanation techniques and argue that the choice depends on the deployment context and the expertise of the user (Barredo Arrieta et al., 2020). Ali et al. (2023) report that tailoring explanations to user types remains an open problem, with most systems defaulting to technically framed outputs. In education, SHAP (Lundberg et al., 2020; Lundberg & Lee, 2017) and LIME (Ribeiro et al., 2016) dominate current applications. Nnadi et al. (2024) found SHAP the most consistent of several methods for attribution stability, Salih et al. (2025) discuss the computational cost of SHAP, its feature independence assumption and the sensitivity of LIME to sampling, and Gunasekara and Saarela (2025) apply both methods to models trained on OULAD. Hooshyar and Yang (2024) report significant variability in the feature importance assigned by SHAP variants and LIME to the same educational model. Critiques

of post-hoc explanation add caution: Rudin (2019) argues that high-stakes decisions should rest on intrinsically interpretable models, and Slack et al. (2020) show that both SHAP and LIME can be manipulated through their perturbation schemes.

Work on the human side of explanation argues that explanations should be contrastive, selective and suited to the recipient (Miller, 2019), and that their effectiveness depends on the fit between content, expertise and decision (Ehsan et al., 2021; Liao & Varshney, 2022). Khosravi et al. (2022) identify several stakeholder groups in educational XAI whose explanation needs differ, and Conati et al. (2021) show that learners with different prior knowledge derive different value from the same explanations. Afzaal et al. (2021) present an early example of SHAP-based explanations delivered within an LMS, but for students rather than for the staff who make intervention decisions.

2.3. Learning analytics dashboards and role-aware design

Learning analytics dashboards (LADs) are the main channel through which analytics reach staff, but reviews report limited theoretical grounding and evaluation (Matcha et al., 2020; Schwendimann et al., 2017). Cabral et al. (2025) identify limited XAI integration as one of the most persistent gaps in AI-powered LADs. Existing designs typically serve one user group: Susnjak et al. (2022) combine descriptive, predictive and prescriptive analytics for learners, Chavan and Mitra (2022) design for instructors, and Kaliisa and Dolonen (2023) found that instructors prioritise summary-level alerts over granular breakdowns. Even within one role, needs vary; in the LADA advising dashboard, senior advisors preferred cohort views while less experienced advisors relied on individual detail (Gutiérrez et al., 2020). Kaliisa et al. (2023) provide a checklist for planning, designing, implementing and evaluating LADs that calls for engagement with the full range of stakeholders. Knowledge engineering offers a structured way to do this: CommonKADS models agents, their tasks and the knowledge they require as distinct artefacts (Schreiber et al., 2000; Studer et al., 1998). Taken together, the reviewed literature shows mature technical components for prediction and explanation, but no staff-facing design that serves instructors, advisors and administrators from one prediction system with role-specific explanations.

3. Method

3.1. Research design

The study follows the Design Science Research Methodology (Hevner et al., 2004; Peffers et al., 2007), with CommonKADS (Schreiber et al., 2000) as the knowledge engineering and design framework applied within it. The Technology Acceptance Model (Davis, 1989) was used only as a design lens for perceived usefulness and ease of use, and the checklist of Kaliisa et al. (2023) guided dashboard design. Table 1 summarises how each of the six CommonKADS models was instantiated. Because no institutional partner or practitioners were involved, the Organisation and Agent Models were populated from published literature rather than from knowledge elicitation, which CommonKADS would ordinarily expect; the resulting role characterisations are therefore well-sourced hypotheses rather than documented accounts.

Table 1. Instantiation of the six CommonKADS models in this study.

Model	Instantiation	Resulting design decision
Organisation	How at-risk information reaches instructors, advisors and administrators and where the flow breaks down (scores without explanation; one view for all roles), from published accounts (Kaliisa et al., 2023; Leitner et al., 2019; Tsai & Gašević, 2017).	Three staff roles served from one shared prediction system; students out of scope.
Task	Risk assessment, explanation generation, role assignment and explanation delivery; intervention left to the human agent (Verbert et al., 2013).	Layered architecture separating data, prediction, explanation, role assignment and presentation.
Agent	Knowledge, capabilities and explanation needs of each role, from literature rather than elicitation (Haque et al., 2023; Kostopoulos et al., 2024).	Three information briefs at individual, caseload and cohort granularity (Table 2).
Knowledge	Engagement, assessment and demographic features and their SHAP attributions (Kuzilek et al., 2017; Lundberg & Lee, 2017).	Feature families with a human-readable feature dictionary.
Communication	Selective explanations phrased for the recipient (Miller, 2019).	Progressive disclosure, side-by-side SHAP and LIME, role-gated views.
Design	Mapping of the five models above onto the technical architecture (Schreiber et al., 2000).	Five-layer architecture with role-based access control (Section 4).

3.2. Dataset and target

OULAD comprises seven relational tables covering registrations, demographics, module presentations, assessments, assessment results, VLE resources and VLE interaction logs (Kuzilek et al., 2017). These were joined into one row per student, module and presentation, giving 32,593 observations. Following the convention in the OULAD literature (Jin et al., 2024), Withdrawn ($n = 10,156$) and Fail ($n = 7,052$) outcomes were mapped to the at-risk class (17,208 observations, 52.8%) and Pass ($n = 12,361$) and Distinction ($n = 3,024$) to the not-at-risk class (15,385, 47.2%). The classes are only mildly imbalanced. Because withdrawals account for 59% of the positive class, the model is best understood as substantially a withdrawal predictor.

3.3. Features

The model input comprised 24 columns (17 numeric and 7 categorical), expanding to 59 features after one-hot encoding (Appendix B). The features form three families. The engagement family contains total VLE clicks, active days, mean daily clicks and clicks in 7-, 14- and 28-day windows before each assessment deadline, following the deadline-anchored strategy of Hlosta et al. (2017); a programmatic check raises an error if any window includes activity on or after its deadline. The assessment family contains the number of assessments submitted, mean and minimum score, lateness, late submissions and banked results. The demographic and registration family contains module, gender, region, highest education, index of multiple deprivation (IMD) band, age band, disability, previous attempts, studied credits, registration date and presentation length. Counts were set to zero where no activity was recorded; other missing values were imputed inside the model pipeline (median or most frequent value) using the training folds only. All behavioural and assessment features were computed over the full course presentation, so the evaluation measures retrospective classification rather than performance at a predefined early-warning point.

3.4. Models and evaluation protocol

Three ensembles representing different points on the bias-variance spectrum were compared: Random Forest (Breiman, 2001), XGBoost (Chen & Guestrin, 2016) and scikit-learn

195 Gradient Boosting (Friedman, 2001; Pedregosa et al., 2011). Evaluation used four-fold
196 GroupKFold cross-validation grouped by course presentation (2013B, 2013J, 2014B and
197 2014J), so no presentation appears in both the training and validation partitions of a fold,
198 avoiding the leakage pattern described by Krstajic et al. (2014). The folds were not stratified,
199 since stratification is incompatible with grouping on a variable with four values. XGBoost
200 used a fixed configuration (300 estimators, maximum depth 5, learning rate 0.05, subsample
201 0.8 and column subsample 0.8) without hyperparameter search or resampling, and the other
202 two models were run under the identical preprocessing pipeline and protocol. Accuracy,
203 precision, recall and F1-score were computed over out-of-fold predictions for all 32,593
204 observations at a 0.5 threshold, together with ROC-AUC. F1-score was used for comparative
205 reporting, while recall was the primary criterion for selecting the prototype model: a student
206 who is at risk but not flagged receives no support, whereas a student flagged unnecessarily is
207 reviewed by a practitioner and can be set aside.

208 *3.5. Explanation analysis*

209 Global importance was computed with TreeSHAP (Lundberg et al., 2020) using the final
210 fitted model across all 32,593 observations, ranking the 59 transformed features by mean
211 absolute SHAP value in log-odds units; these values explain the final model and are not fold-
212 specific. LIME (Ribeiro et al., 2016) was run on a random sample of 40 observations
213 (NumPy default_rng, seed 42), with explanations aggregated by mean absolute weight.
214 Agreement was assessed at two levels: overlap of the aggregate top-10 rankings, and the
215 mean overlap between each sampled instance's top-5 SHAP and top-5 LIME features. To test
216 whether the ordering of two near-duplicate engagement features was stable, the final model
217 was refitted with seven further random seeds. The explanation stability and fidelity checks
218 originally specified, following Visani et al. (2022), were not completed, so the explanation
219 evaluation is limited to cross-method agreement.

220 4. System design and implementation

221 4.1. Architecture

222 The system has five layers (Fig. 1): a Data Layer for feature engineering, a Prediction Engine
223 for risk classification, an Explainability Module for SHAP and LIME, an Interface/API layer
224 enforcing authentication and role-based access control, and a Presentation Layer with three
225 role-differentiated views. The layers communicate through REST endpoints, so each can be
226 revised independently, and explanation is treated as a core architectural component rather
227 than an addition to the model (Mohseni et al., 2021). The fitted preprocessing pipeline and
228 model are saved as one versioned artefact with a manifest recording the training date, feature
229 dictionary version and evaluation metrics, so any prediction returned by the API can be
230 traced to the model version and feature contract that produced it. A versioned feature
231 dictionary also carries a human-readable label for each feature, so explanations are shown
232 with meaningful labels rather than variable names (Miller, 2019).




Fig. 1. Implemented system architecture, showing the data, prediction, explainability, API and presentation layers, the model development environment and the two LMS integration paths tested. All services run locally under Docker Compose.

4.2. Explanation services

The Explainability Module has two services (Fig. 2). A global service computes TreeSHAP values once per model version and caches them for the cohort-level views. A local service computes a local SHAP explanation and an independent LIME explanation on demand for a single prediction, using the exact feature vector the model scored. Caching the global values reflects the higher computing time of SHAP reported by Salih et al. (2025), whereas LIME relies on perturbation and must be regenerated for each instance. Both explanations are shown side by side so that a user can see where they agree and where they differ, following the argument of Salih et al. (2025) for using the two methods together. Following the selectivity property of explanations (Miller, 2019) and evidence that comprehension degrades

when full attribution vectors are displayed (Poursabzi-Sangdeh et al., 2021), the interface shows a small number of leading contributions rather than the complete vector.




Fig. 2. Dual-service explainability pipeline: global TreeSHAP values are pre-computed and cached for cohort views, while LIME and local SHAP are generated on demand for individual students.


4.3. Role-aware views

Each view was designed around an information brief derived from the Agent Model (Table 2). The instructor view gives per-student detail with SHAP and LIME shown side by side, following evidence that teachers act on dashboard information at the level of the individual learner (Molenaar & Knoop-van Campen, 2019). The advisor view averages the reasoning across a caseload of students at or above the watch threshold, following the caseload orientation of LADA (Gutiérrez et al., 2020). The administrator view shows institution-wide model performance and a full-population reasoning trail, in line with institutional decision-support patterns (Tsai & Gašević, 2017). All three views reuse one "reasoning trail" component, a chart of signed SHAP and LIME contributions, at different levels of aggregation, so the roles differ in the granularity rather than the visual form of the knowledge they receive (Sokol & Flach, 2020). Progressive disclosure is applied throughout: compact risk summaries appear by default and attribution detail on request (Springer & Whittaker, 2020). Fig. 3 shows the three views.


266 Table 2. Role-based dashboard design.

Role	Decision	Information needed	Dashboard component	Granularity
Instructor	Whether to act within a course that is currently running.	Class-level engagement in the current period and per-student risk drivers at the level of assessments and weekly activity.	Per-student reasoning trail with SHAP and LIME side by side.	Individual
Advisor	How to approach a particular student and prioritise a caseload.	Views across modules for a caseload, with demographic context suited to one-to-one consultations.	Caseload-averaged reasoning trail for students at or above the watch threshold.	Caseload
Administrator	Where to direct institutional attention and resources.	Cohort-level aggregations and module-over-module comparisons.	Institution-wide performance summary and full-population reasoning trail.	Cohort


267



(a) Instructor view: per-student SHAP and LIME reasoning trail.



(b) Advisor view: caseload-averaged reasoning trail for students at or above the watch threshold.




(c) Administrator view: institution-wide performance summary and full-population reasoning trail.

Fig. 3. The three role-gated dashboard views rendered from the same prediction and explanation services.

276 4.4. Security and LMS integration

277 Role separation is enforced at both the LMS and the API, following the defence-in-depth
278 principle of Saltzer and Schroeder (1975). The API accepts either a JSON Web Token (Jones
279 et al., 2015) from its own login or, for server-to-server calls from the Moodle plugin, a shared
280 service key, and it checks the forwarded role against each endpoint. The visible view is gated
281 by the authenticated session, with no client-side mechanism to switch roles. Two integration
282 paths were tested locally in Docker. In Moodle 5.2 the dashboard was integrated as a native
283 local plugin, with all 22 module presentations loaded as courses, 28,785 students as users and
284 32,593 registrations as enrolments; login, role and course context come from Moodle, and
285 cross-role and cross-course requests were refused when tested (Fig. 4). In Canvas, the
286 dashboard was embedded through an External URL module item, because the available
287 Canvas image predates LTI 1.3 support, so role resolution on that path relies on the
288 dashboard’s own authentication. Both integrations establish technical feasibility rather than
289 deployment readiness.



290
291 Fig. 4. At-risk insights page of the Moodle plugin in course FFF-2014J, showing the SHAP and LIME
292 reasoning trail for one student.

5. Results

5.1. Predictive performance

Table 3 reports out-of-fold performance for the three models. The prototype XGBoost model reached an accuracy of 0.8774, precision of 0.8659, recall of 0.9085, F1-score of 0.8867 and ROC-AUC of 0.9624. Its confusion matrix (Fig. 5) contains 15,634 true positives, 1,574 false negatives, 12,963 true negatives and 2,422 false positives, giving a specificity of 0.8426; about 13% of positive flags were therefore false. Random Forest obtained the highest F1-score (0.8976), accuracy and precision, while XGBoost had the highest recall. The spread in F1-score across the three models was 0.011. Under the recall-first selection criterion, XGBoost was retained for the prototype.

Table 3. Out-of-fold performance under four-fold GroupKFold cross-validation grouped by course presentation (n = 32,593; threshold 0.5).

Model	Accuracy	Precision	Recall	F1-score	Prototype
XGBoost	0.8774	0.8659	0.9085	0.8867	Yes
Random Forest	0.8921	0.8997	0.8955	0.8976	No
Gradient Boosting	0.8869	0.8904	0.8962	0.8932	No

Note. ROC-AUC for XGBoost was 0.9624. Bold marks the best value in each column.




Fig. 5. Out-of-fold accuracy, precision, recall and F1-score of the XGBoost prototype (left) and its confusion matrix at the 0.5 threshold (right), across all 32,593 observations.

5.2. Variability across course presentations

Because each fold holds out one presentation, performance can be reported for each unseen presentation (Table 4, Fig. 6). Accuracy ranged from 0.8456 (2014J) to 0.9161 (2013J), a spread of more than seven percentage points, and the error pattern reversed: the 2013

presentations had precision above 0.95 with recall of 0.86 to 0.88, while the 2014 presentations had precision of 0.81 to 0.82 with recall up to 0.9627. With only four presentation-level folds, the source of this variation cannot be attributed precisely, but the model's operating point was not stable across cohorts even with the model and threshold held fixed.

Table 4. XGBoost out-of-fold performance with each course presentation held out in turn.

Presentation	n	At-risk rate	Accuracy	Precision	Recall	F1-score
2013B	4,684	0.5527	0.9065	0.9506	0.8764	0.9120
2013J	8,845	0.4941	0.9161	0.9690	0.8577	0.9099
2014B	7,804	0.5697	0.8619	0.8243	0.9627	0.8882
2014J	11,260	0.5154	0.8456	0.8075	0.9197	0.8600




Fig. 6. XGBoost out-of-fold accuracy, precision, recall and F1-score for each held-out course presentation.

5.3. Global SHAP explanations

The number of assessments submitted dominated the global TreeSHAP ranking, with a mean absolute SHAP value of 3.1110, more than four times that of the second feature, active days (0.7087) (Table 5, Fig. 7). For the number of assessments submitted, mean score, active days, total clicks and minimum score, the correlation between feature value and SHAP value was strongly negative (-0.69 to -0.96), so higher values pushed predictions away from the at-risk class. The 28-day pre-deadline click window was an exception, with a positive correlation ($+0.8354$), in the opposite direction to total clicks and active days. The engagement variables

are highly collinear (Spearman correlation of 0.98 between the 28-day window and total clicks; 0.95 and 0.92 for the 14- and 7-day windows), and SHAP attributions are affected by collinearity (Salih et al., 2025). This is one possible explanation, but no dependence analysis was carried out, so the sign is reported as an observed property of the fitted model and not as evidence that late-window activity is itself a risk factor.

Table 5. Top 10 features by global TreeSHAP importance and by aggregate LIME weight.

Rank	SHAP feature	Mean SHAP	LIME feature (40 instances)	Mean weight
1	Assessments submitted	3.1110	Assessments submitted	0.3961
2	Active days	0.7087	Module = FFF	0.1397
3	Module = FFF	0.5338	Module = AAA	0.0995
4	Mean score	0.4790	Mean score	0.0640
5	Total clicks	0.3819	Active days	0.0498
6	Clicks 28 days before deadline	0.3111	Presentation length	0.0447
7	Minimum score	0.2741	Clicks 28 days before deadline	0.0236
8	Presentation length	0.2718	Minimum score	0.0183
9	Module = EEE	0.2285	Registration date	0.0073
10	Mean days late	0.1468	Region = Wales	0.0043

Note. SHAP values are in log-odds units, computed with the final fitted model over all 32,593 observations. Seven features appear in both top-10 lists. Mean instance-level top-5 overlap: 0.635 (SD 0.197).




Fig. 7. Top 15 features by global TreeSHAP importance (mean absolute SHAP value), computed with the final fitted model across all 32,593 observations.

5.4. Agreement between SHAP and LIME

Aggregated over the 40 sampled instances, LIME also ranked the number of assessments submitted first (mean absolute weight 0.3961), followed by the FFF and AAA module indicators and mean score (Fig. 8). The two aggregate top-10 rankings shared seven features: assessments submitted, active days, module FFF, mean score, the 28-day click window, minimum score and presentation length. SHAP additionally ranked total clicks, module EEE and mean lateness in its top ten, whereas LIME ranked module AAA, registration date and the Wales region indicator. At the instance level, the mean top-5 overlap was 0.635 (SD 0.197): on average just over three of a student's top five SHAP features also appeared in that student's top five LIME features. LIME tended to rank module indicators higher, consistent with its local, sample-based construction being more sensitive to instances in which module identity is locally decisive. In representative cases, such as a student in FFF-2014J predicted at 0.9998 (true label at risk) and a student in FFF-2014B predicted at 0.0038 (true label not at risk), both methods led with the number of assessments submitted.




Fig. 8. Global SHAP top-10 features (left) and aggregate LIME top-10 features over 40 sampled instances (right).

5.5. Stability of near-duplicate feature rankings

The 28-day click window (0.3111) and total clicks (0.3819) sat close together in the SHAP ranking. Across eight fits with different random seeds, the 28-day window ranked above total clicks in four fits and below it in the other four, while the combined importance of the two features stayed within 0.57 to 0.69. Their relative order is therefore not a stable property of the model: where two features carry almost the same information, the ensemble distributes

credit between them according to its random subsampling. The combined share of pre-deadline and total engagement was stable and well below that of assessment submission. In a separate execution, the out-of-fold metrics were reproduced exactly, matching those recorded in the model manifest, including the confusion matrix cell for cell.

6. Discussion

6.1. Predictive performance under a strict protocol

The figures in Table 3 are lower than the highest accuracies reported on OULAD, such as the 96.9% reported by Wang (2025), but they are not comparable. Grouping folds by presentation means every prediction was made by a model that had not seen the student's cohort, which is closer to applying an existing model to a new presentation than the stratified random splits used in much of the literature (Jin et al., 2024; Krstajic et al., 2014). The per-presentation results reinforce this: one fixed model under one fixed protocol varied by more than seven accuracy points depending only on which presentation was held out. If protocol and cohort produce this much variation within a single evaluation, numerical differences between studies that differ in split, label definition and feature horizon carry little information about relative model quality (Pelánek, 2018).

The error profile is asymmetric: recall exceeded precision at the untuned 0.5 threshold, so the model leans towards flagging. The 1,574 false negatives are students for whom no role would have received a basis to act, while the 2,422 false positives represent staff effort directed at students who did not withdraw or fail. The comparison with Random Forest shows that the choice of selection metric encodes a judgement about which error matters more, and stating that judgement explicitly is preferable to inheriting it from a general-purpose metric. Because the operating point changed between cohorts, an institution could not fully anticipate the error profile for a new cohort, which argues for periodic revalidation and for showing the current operating point to the roles that rely on it.

6.2. What the explanations show

The concentration of attribution on the number of assessments submitted has a benefit and a cost. Assessment submission is a concrete behaviour that an instructor or advisor can act on,

but a reasoning trail dominated by one feature offers less differentiated information than an explanation layer might suggest, and may be read as confirming what is already visible in the gradebook. The direction of that association is also ambiguous, since a student who has decided to withdraw will stop submitting work. The attributions describe how the fitted model reached its output, not what caused a student's situation, and the dashboard and this paper use the language of contribution and association rather than cause.


The partial agreement between the methods is the most distinctive explainability finding. The methods agree strongly in aggregate but are not interchangeable for a single prediction, which is the level at which an instructor or advisor would consult them. This is not an error in either method: TreeSHAP computes exact Shapley values for the tree ensemble, whereas LIME fits a sparse local surrogate to sampled perturbations, so the two answer related but different questions (Lundberg et al., 2020; Ribeiro et al., 2016). The measured overlap converts the principled argument for showing both methods (Salih et al., 2025) into an argument from evidence within this system: showing one method alone would hide a measurable degree of disagreement. Where the methods agree, a practitioner has a firmer basis for treating a contributor as a stable part of the model's reasoning; where they diverge, the appropriate response is more caution about that explanation. The finding is consistent with the variability reported by Hooshyar and Yang (2024) and adds a quantity to it for this model and dataset. The 28-day window result shows why direction must be inspected alongside magnitude: a ranking alone would have suggested that heavier late activity is a warning sign, which the evidence does not support.

6.3. Role-aware design and the limits of the system

The rationale for differentiating views is that the three roles act on different objects. The prediction and its attribution are the same for each role, but an instructor decides whether to change something in a running course, an advisor decides how to approach a particular student, and an administrator decides where to direct institutional attention. A single view would require at least two of them to aggregate or disaggregate information themselves. The technical results also indicate what each role specifically needs. For advisors, a caseload

421 drawn from a cohort resembling 2014B (recall 0.9627, precision 0.8243) would be larger and
 422 contain more students who were not going to withdraw, so a caseload view that does not
 423 show current precision leaves the advisor unable to weigh each flag. For administrators, the
 424 between-cohort variability is arguably the most useful single output, since it indicates how
 425 often the model needs revisiting.

426 CommonKADS contributed traceability from institutional context to interface: the Task
 427 Model's decomposition into risk assessment, explanation generation, role assignment and
 428 explanation delivery is mirrored by the five-layer architecture, and each view's granularity
 429 traces to the Agent Model. The Task Model stops at explanation delivery. Intervention is a
 430 task for the human agent, and the system is designed to support human interpretation without
 431 taking over the intervention decision (Fig. 9). An attribution identifying low assessment
 432 submission narrows where a practitioner should look but does not indicate what to do, since a
 433 student struggling with the material, one managing illness and one who has decided to leave
 434 may present the same feature values. A flag with a plausible explanation may still be a false
 435 positive, as roughly 13% of flags were here, and each would have been accompanied by a
 436 coherent reasoning trail. Whether side-by-side explanations help practitioners calibrate their
 437 reliance or make the interface harder to read can only be settled by observation; Brdnik et al.
 438 (2023) show how perceived trust and satisfaction with SHAP and LIME can be measured in a
 439 learning analytics setting.



440
 441 Fig. 9. Relationship between prediction, explanation, human interpretation and intervention. The
 442 system performs the first two stages; interpretation and intervention remain with the practitioner.

443 6.4. *Candidate design guidelines*

444 The findings suggest six guidelines for explainable learning analytics dashboards, offered as
445 candidates for empirical validation:

- 446 • Report the agreement between explanation methods where more than one is available,
447 since a single method presents a more settled picture of local model behaviour than the
448 attributions support.
- 449 • Let users inspect the direction of each contribution, not only its magnitude, and check
450 for collinearity artefacts before attributions reach practitioners.
- 451 • Use the language of contribution and association rather than causal or prescriptive
452 language.
- 453 • Show the current operating point (the balance between missed students and false
454 alarms) to the roles that rely on it, and revalidate it for each new cohort.
- 455 • Keep the boundary between explanation delivery by the system and intervention by the
456 human, rather than extending the system to recommend actions its explanation methods
457 cannot justify.
- 458 • Report subgroup performance before the system is used to direct resources, and enforce
459 role-based access to per-student risk information at more than one layer.

460 6.5. *Transferability*

461 All results come from one distance-learning institution, where the LMS is the primary site of
462 study and withdrawal is common, and from presentations in 2013 and 2014. Given that
463 performance moved by more than seven accuracy points across four presentations of the
464 same modules, a model trained at one institution should be expected to need retraining and
465 revalidation elsewhere rather than transfer as a fitted artefact. The CommonKADS design
466 process, the layered architecture, the cached-global and on-demand-local explanation services
467 and the reasoning trail rendered at several levels of aggregation do not depend on OULAD
468 and are the parts most likely to transfer. Transfer would require LMS interaction and
469 assessment data of sufficient completeness, reliably recorded deadlines, an outcome
470 definition suited to the institution, roles that correspond to those modelled here and a

governance position on the use of per-student risk predictions. Where teaching is largely in person and the LMS holds only supplementary material, the engagement features that carry much of the signal here would be sparse, and the approach would need to be reconsidered at the level of the Knowledge Model.

6.6. *Limitations*

Five limitations constrain these findings. First, no practitioner used the system, so trust, comprehension, usability and decision quality were not measured, and the role-specific design rests on published accounts rather than elicitation. Second, the features were computed over the full presentation, so the evaluation establishes retrospective classification rather than early-warning performance. Third, calibration of the displayed probabilities, subgroup fairness, explanation stability and fidelity, hyperparameter tuning, class-imbalance handling and trajectory features were planned but not completed, and the LIME analysis rests on 40 instances, with explanations computed in a one-hot feature space that splits categorical attributes across several indicators. Fourth, one registration feature (`withdrew_before_start`) partly encodes the outcome, since all 3,097 observations with this flag are labelled at risk; its mean absolute SHAP value in the final model is 0.00, because the same students have no submitted assessments, but it should be removed in future iterations. Fifth, the LMS integrations were tested in local containers rather than institutional installations.

6.7. *Future work*

Three priorities follow directly from these limitations: an early-warning evaluation with features restricted to the information available at fixed points such as weeks 2, 4, 6 and 8 (Adnan et al., 2021); a calibration study comparing isotonic regression and Platt scaling (Guo et al., 2017; Niculescu-Mizil & Caruana, 2005); and a subgroup fairness assessment across gender, age band, disability and IMD band (Hardt et al., 2016; Mitchell et al., 2019). Beyond these, a study with instructors, advisors and administrators should test how the role-specific views are read and used, drawing on the perceived usefulness and ease of use constructs of TAM, and the approach should be applied to data from other institutions and modes of study.

498 **7. Conclusion**

499 This study designed, implemented and technically evaluated a role-aware explainable AI
500 dashboard that delivers predictions and explanations from one model to instructors, advisors
501 and administrators at different levels of aggregation. On RQ1, a cached global TreeSHAP
502 service paired with on-demand local SHAP and LIME ran at full dataset scale and produced
503 attributions traceable to a versioned model and feature contract. On RQ2, the CommonKADS
504 models yielded a traceable specification of three role-gated views that share one reasoning
505 trail at individual, caseload and cohort granularity. On RQ3, the technical preconditions for
506 supporting practitioner interpretation are in place: the model discriminates well on unseen
507 presentations, and its explanations are human-readable and cross-checked, with a measured
508 instance-level SHAP and LIME overlap of 0.635 that supports showing both methods.
509 Results describe retrospective classification on a single dataset, and whether the dashboard
510 improves practitioners' trust, understanding or decisions remains to be tested. The prototype
511 provides a verified technical foundation for that evaluation.

512 **CRediT authorship contribution statement**

513 Daniel Larry: Conceptualization, Methodology, Software, Formal analysis, Investigation,
514 Data curation, Visualization, Writing – original draft. Zainab Shamsiya Usman: Supervision,
515 Methodology, Writing – review & editing. [Authors to confirm.]

516 **Declaration of competing interest**

517 The authors declare that they have no known competing financial interests or personal
518 relationships that could have appeared to influence the work reported in this paper. [Authors
519 to confirm.]

520 **Funding**

521 This research did not receive any specific grant from funding agencies in the public,
522 commercial, or not-for-profit sectors. [Authors to confirm.]

523 **Ethics statement**

524 The study used only the publicly available, anonymised Open University Learning Analytics
525 Dataset. No primary data were collected from students, instructors, advisors or
526 administrators.

527 **Data availability**

528 OULAD is publicly available (Kuzilek et al., 2017). The prototype source code, including
529 data preparation, model training, explanation services, the Flask API, the React dashboard
530 and the Moodle plugin, is available at [https://github.com/daniel-larry/role-aware-xai-](https://github.com/daniel-larry/role-aware-xai-dashboard)
531 dashboard.

532 **Declaration of generative AI and AI-assisted technologies in the writing process**

533 During the preparation of this work the authors used Claude (Anthropic) to condense the
534 master's thesis on which this article is based into journal format and to verify and format the
535 references. After using this tool, the authors reviewed and edited the content as needed and
536 take full responsibility for the content of the published article. [Authors to confirm.]

537 **Acknowledgements**

538 This article is based on the first author's MSc thesis in Computer Science at the American
539 University of Nigeria.

540 **Appendix A. Reproducibility details**

541 Table A.1. Reproducibility details of the evaluated pipeline.

Item	Value
Software	Python 3.11.16; scikit-learn 1.9.1; xgboost 3.2.0; pandas 3.0.6; NumPy 2.4.6; SHAP 0.51.0; LIME 0.2.0.1; joblib 1.6.0; Flask 3.1.3; PyJWT 2.15.1; gunicorn 26.2.0.
Hardware	Docker container on a Linux x86_64 host with 4 virtual CPUs and 16 GB of memory; no GPU.
Code version	github.com/daniel-larry/role-aware-xai-dashboard , branch main, commit 7c0068b.
Model artefact	Serialised joblib pipeline (preprocessing and model) with a JSON manifest; identifier xgb-20260929-112741.
Model configuration	XGBoost, 300 estimators, maximum depth 5, learning rate 0.05, subsample 0.8, column subsample 0.8; threshold 0.5; no resampling or hyperparameter search.

Item	Value
Validation	GroupKFold, four folds grouped by code_presentation (2013B, 2013J, 2014B, 2014J); out-of-fold metrics over all 32,593 observations.
Random seeds	random_state = 42 throughout; LIME sample of 40 instances with NumPy default_rng, seed 42; seed-stability refits with random_state 1 to 7.

Appendix B. Feature definitions

Features marked "Full presentation" were computed over the whole course presentation and are not available at an early-warning point as constructed.

Table B.1. Input features of the evaluated model.

Feature	Definition	Availability
num_of_prev_attempts	Number of previous attempts at the module	Registration
studied_credits	Total credits of the modules being studied	Registration
date_registration	Registration date relative to module start (days)	Registration
withdrew_before_start	Unregistered on or before the module start (see Section 6.6)	Before start
module_presentation_length	Length of the module presentation (days)	Before start
n_assessments_submitted	Number of assessments submitted	Full presentation
mean_score	Mean score of submitted assessments	Full presentation
min_score	Lowest score of submitted assessments	Full presentation
mean_days_late	Mean lateness of submissions (on-time counted as 0)	Full presentation
n_late_submissions	Number of submissions after the deadline	Full presentation
n_banked	Assessment results carried over from a previous presentation	Full presentation
clicks_7d/14d/28d_before_deadline	VLE clicks in the 7, 14 and 28 days before assessment deadlines	Full presentation
total_clicks	Total VLE clicks	Full presentation
active_days	Number of days with VLE activity	Full presentation
mean_daily_clicks	Total clicks divided by active days	Full presentation
code_module, gender, region, highest_education, imd_band, age_band, disability	Categorical attributes, one-hot encoded	Registration

549 **References**

- 550 Adnan, M., Habib, A., Ashraf, J., Mussadiq, S., Raza, A. A., Abid, M., Bashir, M., & Khan,
551 S. U. (2021). Predicting at-risk students at different percentages of course length for
552 early intervention using machine learning models. *IEEE Access*, 9, 7519–7539.
553 <https://doi.org/10.1109/ACCESS.2021.3049446>
- 554 Afzaal, M., Nouri, J., Zia, A., Papapetrou, P., Fors, U., Wu, Y., Li, X., & Weegar, R. (2021).
555 Explainable AI for data-driven feedback and intelligent action recommendations to
556 support students self-regulation. *Frontiers in Artificial Intelligence*, 4, Article 723447.
557 <https://doi.org/10.3389/frai.2021.723447>
- 558 Ahmed, W., Wani, M. A., Plawiak, P., Meshoul, S., Mahmoud, A., & Hammad, M. (2025).
559 Machine learning-based academic performance prediction with explainability for
560 enhanced decision-making in educational institutions. *Scientific Reports*, 15(1),
561 Article 26879. <https://doi.org/10.1038/s41598-025-12353-4>
- 562 Ali, S., Abuhmed, T., El-Sappagh, S., Muhammad, K., Alonso-Moral, J. M., Confalonieri, R.,
563 Guidotti, R., Del Ser, J., Díaz-Rodríguez, N., & Herrera, F. (2023). Explainable
564 Artificial Intelligence (XAI): What we know and what is left to attain Trustworthy
565 Artificial Intelligence. *Information Fusion*, 99, Article 101805.
566 <https://doi.org/10.1016/j.inffus.2023.101805>
- 567 Baker, R. S., & Hawn, A. (2022). Algorithmic bias in education. *International Journal of*
568 *Artificial Intelligence in Education*, 32(4), 1052–1092.
569 <https://doi.org/10.1007/s40593-021-00285-9>
- 570 Barredo Arrieta, A., Díaz-Rodríguez, N., Del Ser, J., Bennetot, A., Tabik, S., Barbado, A.,
571 García, S., Gil-López, S., Molina, D., Benjamins, R., Chatila, R., & Herrera, F.
572 (2020). Explainable Artificial Intelligence (XAI): Concepts, taxonomies,
573 opportunities and challenges toward responsible AI. *Information Fusion*, 58, 82–115.
574 <https://doi.org/10.1016/j.inffus.2019.12.012>
- 575 Brdnik, S., Podgorelec, V., & Šumak, B. (2023). Assessing perceived trust and satisfaction
576 with multiple explanation techniques in XAI-enhanced learning analytics. *Electronics*,
577 12(12), 2594. <https://doi.org/10.3390/electronics12122594>

578 Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32.
579 <https://doi.org/10.1023/A:1010933404324>

580 Cabral, L., Pinto, R., & Gonçalves, G. (2025). AI-powered learning analytics dashboards: A
581 systematic review of applications, techniques, and research gaps. *Discover Education*,
582 4(1), Article 525. <https://doi.org/10.1007/s44217-025-00964-y>

583 Chavan, P., & Mitra, R. (2022). Tcherly: *Journal of Learning Analytics*, 9(3), 125–151.
584 <https://doi.org/10.18608/jla.2022.7555>

585 Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of*
586 *the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data*
587 *Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>

588 Conati, C., Barral, O., Putnam, V., & Rieger, L. (2021). Toward personalized XAI: A case
589 study in intelligent tutoring systems. *Artificial Intelligence*, 298, Article 103503.
590 <https://doi.org/10.1016/j.artint.2021.103503>

591 Davis, F. D. (1989). Perceived usefulness, perceived ease of use, and user acceptance of
592 information technology. *MIS Quarterly*, 13(3), 319–340.
593 <https://doi.org/10.2307/249008>

594 Ehsan, U., Liao, Q. V., Muller, M., Riedl, M. O., & Weisz, J. D. (2021). Expanding
595 explainability: Towards social transparency in AI systems. *Proceedings of the 2021*
596 *CHI Conference on Human Factors in Computing Systems*, 1–19.
597 <https://doi.org/10.1145/3411764.3445188>

598 Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *The*
599 *Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>

600 Guevara-Reyes, R., Ortiz-Garcés, I., Andrade, R., Cox-Riquetti, F., & Villegas-Ch, W.
601 (2025). Machine learning models for academic performance prediction:
602 Interpretability and application in educational decision-making. *Frontiers in*
603 *Education*, 10, Article 1632315. <https://doi.org/10.3389/feduc.2025.1632315>

604 Gunasekara, S., & Saarela, M. (2025). Explainable AI in education: Techniques and
605 qualitative assessment. *Applied Sciences*, 15(3), 1239.
606 <https://doi.org/10.3390/app15031239>

607 Guo, C., Pleiss, G., Sun, Y., & Weinberger, K. Q. (2017). On calibration of modern neural
608 networks. *Proceedings of the 34th International Conference on Machine Learning,*
609 *Proceedings of Machine Learning Research, 70*, 1321–1330.
610 <https://proceedings.mlr.press/v70/guo17a.html>

611 Gutiérrez, F., Seipp, K., Ochoa, X., Chiluíza, K., De Laet, T., & Verbert, K. (2020). LADA:
612 A learning analytics dashboard for academic advising. *Computers in Human*
613 *Behavior, 107*, Article 105826. <https://doi.org/10.1016/j.chb.2018.12.004>

614 Haque, A. B., Islam, A. K. M. N., & Mikalef, P. (2023). Explainable artificial intelligence
615 (XAI) from a user perspective: A synthesis of prior literature and problematizing
616 avenues for future research. *Technological Forecasting and Social Change, 186*,
617 Article 122120. <https://doi.org/10.1016/j.techfore.2022.122120>

618 Hardt, M., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning.
619 *Advances in Neural Information Processing Systems, 29*, 3315–3323.
620 <https://proceedings.neurips.cc/paper/2016/hash/6a9659feb1216f14f7384ba499518b38>
621 -Abstract.html

622 Hevner, A. R., March, S. T., Park, J., & Ram, S. (2004). Design science in information
623 systems research. *MIS Quarterly, 28*(1), 75–106. <https://doi.org/10.2307/25148625>

624 Hlosta, M., Zdrahal, Z., & Zendulka, J. (2017). Ouroboros: Early identification of at-risk
625 students without models based on legacy data. *Proceedings of the Seventh*
626 *International Learning Analytics & Knowledge Conference*, 6–15.
627 <https://doi.org/10.1145/3027385.3027449>

628 Hooshyar, D., & Yang, Y. (2024). Problems with SHAP and LIME in interpretable AI for
629 education: A comparative study of post-hoc explanations and neural-symbolic rule
630 extraction. *IEEE Access, 12*, 137472–137490.
631 <https://doi.org/10.1109/ACCESS.2024.3463948>

632 Hussain, M., Zhu, W., Zhang, W., & Abidi, S. M. R. (2018). Student engagement predictions
633 in an e-learning system and their impact on student course assessment scores.
634 *Computational Intelligence and Neuroscience, 2018*, Article 6347186.
635 <https://doi.org/10.1155/2018/6347186>

636 Jin, L., Wang, Y., Song, H., & So, H.-J. (2024). Predictive modelling with the Open
637 University Learning Analytics Dataset (OULAD): A systematic literature review. In
638 A. M. Olney, I.-A. Chounta, Z. Liu, O. C. Santos, & I. I. Bittencourt (Eds.), *Artificial*
639 *intelligence in education. Posters and late breaking results, workshops and tutorials,*
640 *industry and innovation tracks, practitioners, doctoral consortium and blue sky* (pp.
641 477–484). Springer. https://doi.org/10.1007/978-3-031-64315-6_46

642 Jones, M., Bradley, J., & Sakimura, N. (2015). *JSON Web Token (JWT)* [RFC 7519]. Internet
643 Engineering Task Force. <https://doi.org/10.17487/RFC7519>

644 Kaliisa, R., & Dolonen, J. A. (2023). CADA: A teacher-facing learning analytics dashboard
645 to foster teachers’ awareness of students’ participation and discourse patterns in
646 online discussions. *Technology, Knowledge and Learning*, 28(3), 937–958.
647 <https://doi.org/10.1007/s10758-022-09598-7>

648 Kaliisa, R., Jivet, I., & Prinsloo, P. (2023). A checklist to guide the planning, designing,
649 implementation, and evaluation of learning analytics dashboards. *International*
650 *Journal of Educational Technology in Higher Education*, 20(1), Article 28.
651 <https://doi.org/10.1186/s41239-023-00394-6>

652 Khosravi, H., Buckingham Shum, S., Chen, G., Conati, C., Tsai, Y.-S., Kay, J., Knight, S.,
653 Martinez-Maldonado, R., Sadiq, S., & Gašević, D. (2022). Explainable artificial
654 intelligence in education. *Computers and Education: Artificial Intelligence*, 3, Article
655 100074. <https://doi.org/10.1016/j.caeai.2022.100074>

656 Kostopoulos, G., Davrazos, G., & Kotsiantis, S. (2024). Explainable artificial intelligence-
657 based decision support systems: A recent review. *Electronics*, 13(14), 2842.
658 <https://doi.org/10.3390/electronics13142842>

659 Krstajic, D., Buturovic, L. J., Leahy, D. E., & Thomas, S. (2014). Cross-validation pitfalls
660 when selecting and assessing regression and classification models. *Journal of*
661 *Cheminformatics*, 6(1), Article 10. <https://doi.org/10.1186/1758-2946-6-10>

662 Kuzilek, J., Hlostá, M., & Zdrahal, Z. (2017). Open University Learning Analytics dataset.
663 *Scientific Data*, 4(1), Article 170171. <https://doi.org/10.1038/sdata.2017.171>

664 Leitner, P., Ebner, M., & Ebner, M. (2019). Learning analytics challenges to overcome in
 665 higher education institutions. In D. Ifenthaler, D.-K. Mah, & J. Y.-K. Yau (Eds.),
 666 *Utilizing learning analytics to support study success* (pp. 91–104). Springer.
 667 https://doi.org/10.1007/978-3-319-64792-0_6

668 Liao, Q. V., & Varshney, K. R. (2022). *Human-centered explainable AI (XAI): From*
 669 *algorithms to user experiences* (arXiv:2110.10790). arXiv.
 670 <https://doi.org/10.48550/arXiv.2110.10790>

671 Lundberg, S. M., Erion, G., Chen, H., DeGrave, A., Prutkin, J. M., Nair, B., Katz, R.,
 672 Himmelfarb, J., Bansal, N., & Lee, S.-I. (2020). From local explanations to global
 673 understanding with explainable AI for trees. *Nature Machine Intelligence*, 2(1), 56–
 674 67. <https://doi.org/10.1038/s42256-019-0138-9>

675 Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions.
 676 *Advances in Neural Information Processing Systems*, 30, 4765–4774.
 677 [https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b6776](https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html)
 678 [7-Abstract.html](https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html)

679 Matcha, W., Uzir, N. A., Gašević, D., & Pardo, A. (2020). A systematic review of empirical
 680 studies on learning analytics dashboards: A self-regulated learning perspective. *IEEE*
 681 *Transactions on Learning Technologies*, 13(2), 226–245.
 682 <https://doi.org/10.1109/TLT.2019.2916802>

683 Miller, T. (2019). Explanation in artificial intelligence: Insights from the social sciences.
 684 *Artificial Intelligence*, 267, 1–38. <https://doi.org/10.1016/j.artint.2018.07.007>

685 Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E.,
 686 Raji, I. D., & Gebru, T. (2019). Model cards for model reporting. *Proceedings of the*
 687 *Conference on Fairness, Accountability, and Transparency*, 220–229.
 688 <https://doi.org/10.1145/3287560.3287596>

689 Mohseni, S., Zarei, N., & Ragan, E. D. (2021). A multidisciplinary survey and framework for
 690 design and evaluation of explainable AI systems. *ACM Transactions on Interactive*
 691 *Intelligent Systems*, 11(3–4), Article 24. <https://doi.org/10.1145/3387166>

692 Molenaar, I., & Knoop-van Campen, C. A. N. (2019). How teachers make dashboard
693 information actionable. *IEEE Transactions on Learning Technologies*, 12(3), 347–
694 355. <https://doi.org/10.1109/TLT.2018.2851585>

695 Niculescu-Mizil, A., & Caruana, R. (2005). Predicting good probabilities with supervised
696 learning. *Proceedings of the 22nd International Conference on Machine Learning*,
697 625–632. <https://doi.org/10.1145/1102351.1102430>

698 Nnadi, L. C., Watanobe, Y., Rahman, Md. M., & John-Otumu, A. M. (2024). Prediction of
699 students' adaptability using explainable AI in educational machine learning models.
700 *Applied Sciences*, 14(12), 5141. <https://doi.org/10.3390/app14125141>

701 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M.,
702 Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D.,
703 Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine learning in
704 Python. *Journal of Machine Learning Research*, 12, 2825–2830.
705 <https://jmlr.org/papers/v12/pedregosa11a.html>

706 Peffers, K., Tuunanen, T., Rothenberger, M. A., & Chatterjee, S. (2007). A design science
707 research methodology for information systems research. *Journal of Management*
708 *Information Systems*, 24(3), 45–77. <https://doi.org/10.2753/MIS0742-1222240302>

709 Pelánek, R. (2018). The details matter: Methodological nuances in the evaluation of student
710 models. *User Modeling and User-Adapted Interaction*, 28(3), 207–235.
711 <https://doi.org/10.1007/s11257-018-9204-y>

712 Poursabzi-Sangdeh, F., Goldstein, D. G., Hofman, J. M., Wortman Vaughan, J., & Wallach,
713 H. (2021). Manipulating and measuring model interpretability. *Proceedings of the*
714 *2021 CHI Conference on Human Factors in Computing Systems*, Article 237.
715 <https://doi.org/10.1145/3411764.3445315>

716 Rebelo Marcolino, M., Reis Porto, T., Thompsen Primo, T., Targino, R., Ramos, V., Marques
717 Queiroga, E., Munoz, R., & Cechinel, C. (2025). Student dropout prediction through
718 machine learning optimization: Insights from Moodle log data. *Scientific Reports*,
719 15(1), Article 9840. <https://doi.org/10.1038/s41598-025-93918-1>

720 Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). “Why should I trust you?”: Explaining the
 721 predictions of any classifier. *Proceedings of the 22nd ACM SIGKDD International*
 722 *Conference on Knowledge Discovery and Data Mining*, 1135–1144.
 723 <https://doi.org/10.1145/2939672.2939778>

724 Rudin, C. (2019). Stop explaining black box machine learning models for high stakes
 725 decisions and use interpretable models instead. *Nature Machine Intelligence*, 1(5),
 726 206–215. <https://doi.org/10.1038/s42256-019-0048-x>

727 Salih, A. M., Raisi-Estabragh, Z., Boscolo Galazzo, I., Radeva, P., Petersen, S. E., Lekadir,
 728 K., & Menegaz, G. (2025). A perspective on explainable artificial intelligence
 729 methods: SHAP and LIME. *Advanced Intelligent Systems*, 7(1), Article 2400304.
 730 <https://doi.org/10.1002/aisy.202400304>

731 Saltzer, J. H., & Schroeder, M. D. (1975). The protection of information in computer
 732 systems. *Proceedings of the IEEE*, 63(9), 1278–1308.
 733 <https://doi.org/10.1109/PROC.1975.9939>

734 Schreiber, G., Akkermans, H., Anjewierden, A., de Hoog, R., Shadbolt, N. R., Van de Velde,
 735 W., & Wielinga, B. J. (2000). *Knowledge engineering and management: The*
 736 *CommonKADS methodology*. MIT Press.
 737 <https://doi.org/10.7551/mitpress/4073.001.0001>

738 Schwendimann, B. A., Rodríguez-Triana, M. J., Vozniuk, A., Prieto, L. P., Boroujeni, M. S.,
 739 Holzer, A., Gillet, D., & Dillenbourg, P. (2017). Perceiving learning at a glance: A
 740 systematic literature review of learning dashboard research. *IEEE Transactions on*
 741 *Learning Technologies*, 10(1), 30–41. <https://doi.org/10.1109/TLT.2016.2599522>

742 Siemens, G., & Long, P. (2011). Penetrating the fog: Analytics in learning and education.
 743 *EDUCAUSE Review*, 46(5), 30–40.
 744 [https://er.educause.edu/articles/2011/9/penetrating-the-fog-analytics-in-learning-and-](https://er.educause.edu/articles/2011/9/penetrating-the-fog-analytics-in-learning-and-education)
 745 [education](https://er.educause.edu/articles/2011/9/penetrating-the-fog-analytics-in-learning-and-education)

746 Slack, D., Hilgard, S., Jia, E., Singh, S., & Lakkaraju, H. (2020). Fooling LIME and SHAP:
 747 Adversarial attacks on post hoc explanation methods. *Proceedings of the AAAI/ACM*

748 *Conference on AI, Ethics, and Society*, 180–186.
749 <https://doi.org/10.1145/3375627.3375830>

750 Sokol, K., & Flach, P. (2020). Explainability fact sheets: A framework for systematic
751 assessment of explainable approaches. *Proceedings of the 2020 Conference on*
752 *Fairness, Accountability, and Transparency*, 56–67.
753 <https://doi.org/10.1145/3351095.3372870>

754 Springer, A., & Whittaker, S. (2020). Progressive disclosure: When, why, and how do users
755 want algorithmic transparency information? *ACM Transactions on Interactive*
756 *Intelligent Systems*, 10(4), 1–32. <https://doi.org/10.1145/3374218>

757 Studer, R., Benjamins, V. R., & Fensel, D. (1998). Knowledge engineering: Principles and
758 methods. *Data & Knowledge Engineering*, 25(1–2), 161–197.
759 [https://doi.org/10.1016/S0169-023X\(97\)00056-6](https://doi.org/10.1016/S0169-023X(97)00056-6)

760 Susnjak, T., Ramaswami, G. S., & Mathrani, A. (2022). Learning analytics dashboard: A tool
761 for providing actionable insights to learners. *International Journal of Educational*
762 *Technology in Higher Education*, 19(1), Article 12. [https://doi.org/10.1186/s41239-](https://doi.org/10.1186/s41239-021-00313-7)
763 [021-00313-7](https://doi.org/10.1186/s41239-021-00313-7)

764 Tsai, Y.-S., & Gašević, D. (2017). Learning analytics in higher education: Challenges and
765 policies. *Proceedings of the Seventh International Learning Analytics & Knowledge*
766 *Conference*, 233–242. <https://doi.org/10.1145/3027385.3027400>

767 Verbert, K., Duval, E., Klerkx, J., Govaerts, S., & Santos, J. L. (2013). Learning analytics
768 dashboard applications. *American Behavioral Scientist*, 57(10), 1500–1509.
769 <https://doi.org/10.1177/0002764213479363>

770 Visani, G., Bagli, E., Chesani, F., Poluzzi, A., & Capuzzo, D. (2022). Statistical stability
771 indices for LIME: Obtaining reliable explanations for machine learning models.
772 *Journal of the Operational Research Society*, 73(1), 91–101.
773 <https://doi.org/10.1080/01605682.2020.1865846>

774 Wang, Y. (2025). Artificial intelligence in student management systems to enhance academic
775 performance monitoring and intervention. *Scientific Reports*, 15(1), Article 35122.
776 <https://doi.org/10.1038/s41598-025-19159-4>

777 Yu, R., Li, Q., Fischer, C., Doroudi, S., & Xu, D. (2020). Towards accurate and fair
778 prediction of college success: Evaluating different sources of student data. In A. N.
779 Rafferty, J. Whitehill, V. Cavalli-Sforza, & C. Romero (Eds.), *Proceedings of the*
780 *13th International Conference on Educational Data Mining* (pp. 292–301).
781 International Educational Data Mining Society.